
SAMPLING FROM THE BAYESIAN POSTER OF AN AGENT-BASED MODEL GIVEN PARTIAL OBSERVATIONS

Daniel Tang
Leeds Institute for Data Analytics*
University of Leeds
Leeds, UK
D.Tang@leeds.ac.uk

February 4, 2021

ABSTRACT

abstract

Keywords Data assimilation, Bayesian inference, Agent based model, Integer linear programming, predator prey model

1 Introduction

Suppose we have a timestepping ABM where agents have a finite number of possible internal states and a finite number of ways of acting on their world. Given this, we can define an ABM as:

- A domain of agent actions $\mathcal{A} = \{a_0 \dots a_n\}$
- A domain of agent states $\mathcal{S} = \{\sigma_0 \dots \sigma_m\}$
- A computer program, $\pi(\psi, \Psi)$, where $\psi \in \mathcal{S}$ is the agent's state and Ψ is a (sparse) vector whose i^{th} element is the number of agents in state σ_i at the start of the timestep. The program can make calls to a random number generator that returns a value $0 \leq r < 1$ with uniform probability, so its returned value may be non-deterministic and we write $P(\pi(\psi, \Psi) = a)$ to be the probability that the program will return value a , given inputs ψ and Ψ . The program returns an action $a \in \mathcal{A}$ which is the action performed by the agent in this timestep.
- An agent transition tensor $F_\phi^{\psi a}$ whose elements give the number of agents in state ϕ that result from an agent in state ψ performing act a (including the acting agent).

As a very simple demonstration, we'll take a spatial predator-prey model where agents are either predator or prey moving about a 2-dimensional grid. In this case, the domain of actions are $\mathbb{A} = \{\text{move-left, move-right, move-up, move-down, reproduce up, reproduce down, reproduce left, reproduce right, be eaten, die}\}$. So the domain of agent states, \mathcal{S} , can be thought of as the class

```
class Agent {  
    Boolean isaPredator  
    Integer xPosition  
    Integer yPosition  
}
```

We can assign an ordering of agent states by placing agent a in the $(a.\text{isaPredator} * G^2 + a.\text{yPosition} * G + a.\text{xPosition})^{th}$ position, where G is the size of the grid.

The timestep function for a predator/prey agent is shown in algorithm 1 (where we assume the grid has periodic boundary conditions so we don't need to worry about its edges). The transition tensor encodes the results of the various

*This project has received funding from the European Research Council (ERC) under the European Union's Horizon 2020 research and innovation programme (grant agreement No. 757455)

Algorithm 1 Timestep of a predator/prey agent

```

function TIMESTEP(agent, otherAgents)
    if agent.isaPredator then
        if RANDOM( ) <  $\beta_0$  and number of prey on neighbouring gridsquares of otherAgents > 0 then
            return reproduce left, right, up or down based on a call to RANDOM( )
        end if
        if RANDOM( ) <  $\gamma_0$  then
            return die
        end if
    else
        if RANDOM( ) <  $\beta_1$  then
            return reproduce left, right, up or down based on a call to RANDOM( )
        end if
        if RANDOM( ) <  $\gamma_1$  then
            return die
        end if
        if RANDOM( ) <  $\delta$  and number of predators on this or neighbouring gridsquares of otherAgents > 0 then
            return be eaten
        end if
    end if
    return move left, right, up or down based on a call to RANDOM( )
end function
    
```

actions. So, for example, the zeroth state, σ_0 , corresponds to a prey at position (0,0). If this agent moves right (i.e. performs action a_1) then the result is a prey at position (1,0) (i.e. in state σ_1) so $F_1^{01} = 1$ etc.

Let a model timestep consist of a matrix E whose elements $e_{\psi a}$ are the number of agents in state ψ that perform act a in this timestep. Similarly, let a model trajectory be a tensor $T_{\psi a}^t$ whose elements are the number of agents in state ψ that perform act a in timestep t .

A trajectory must satisfy a number of constraints in order to be a possible trajectory of an ABM. Since the elements of a trajectory are counts, they must all be non-negative integers, we'll call this the *non-negative integer constraint*

$$\forall t, \psi, a : T_{\psi a}^t \in \mathbb{Z}_{\geq 0} \quad (1)$$

A trajectory, $T_{\psi a}^t$, must also be *continuous* which means that each agent that results from the actions of an agent in timestep t must appear in the actions of timestep $t + 1$. This can be expressed in terms of the *continuity constraints*:

$$\forall t \in 1 \dots n : \forall \phi : \sum_{\psi, a} F_{\phi}^{\psi a} T_{\psi a}^{t-1} - \sum_a T_{\phi a}^t = 0 \quad (2)$$

If a trajectory satisfies constraints 1 and 2 we say that it is *valid*.

Suppose we have a prior belief, $P_0(\Psi)$, over the state of the ABM at time $t = 0$ (in the example of the predator/prey model, let's suppose we know the probability of finding a predator or prey at each gridpoint at $t = 0$, but don't know the exact positions of each agent). Given this, the prior probability of a trajectory is

$$P(T_{\psi a}^t) = \begin{cases} P_0(\sum_a T_{\psi a}^0) \prod_{t, \psi, a} P(\pi(\psi, \sum_b T_{\phi b}^t) = a)^{T_{\psi a}^t} & \text{if } T_{\psi a}^t \text{ is valid} \\ 0 & \text{otherwise} \end{cases}$$

Suppose we also have a set of noisy, aggregate observations, Ω , that have a likelihood function $P(\Omega | T_{\psi a}^t)$. In our running example, let's suppose we periodically note down when and where we see footprints of predator and prey. Our aim is to generate a number of samples from the posterior distribution

$$P(T_{\psi a}^t | \Omega) \propto P(\Omega | T_{\psi a}^t) P(T_{\psi a}^t)$$

We assume that observations are independent given the trajectory and that each individual observation in Ω can be simulated by a computer program $\omega(T_{\psi a}^t)$ which may make calls to a random number generator. The probability, $P(\omega(T_{\psi a}^t) = v)$, that the program returns a value v defines the likelihood of observing value v upon making observation ω on trajectory $T_{\psi a}^t$. i.e. $\omega(T_{\psi a}^t)$ is just a simulation of the observation, which is usually quite easy to express as a computer program and to compute.

In the case of the predator prey model, let's suppose that at the end of every day we go out to a fixed set of inspection sites and check for footprints. Suppose that if predators were present in the grid-square containing an inspection site on

Algorithm 2 Function for observing predator/prey footprints at a site

```

function OBSERVEFOOTPRINTS(T)                                ▷ T is the ABM trajectory
    t ← time of observation
    x ← x-position of observation
    y ← y-position of observation
    G ← size of grid
    ψ ← G2 + Gy + x                                           ▷ state of predator in this gridsquare
    ϕ ← Gy + x                                                  ▷ state of prey in this gridsquare
    Fpred ← FALSE                                              ▷ did we observe predator footprints?
    Fprey ← FALSE                                              ▷ did we observe prey footprints?
    if ∑a Tψat > 0 and RANDOM() < 0.5 then
        Fpred ← TRUE
    end if
    if ∑a Tϕat > 0 and RANDOM() < 0.25 then
        Fprey ← TRUE
    end if
    return Fpred, Fprey
end function
    
```

that day, there's a 50% chance we'll see a footprint, and if prey were present, there's a 25% chance. The observation function for a site is shown in algorithm 2

Given this

$$P(\Omega | T_{\psi a}^t) = \prod_{(\omega, v) \in \Omega} P(\omega(T_{\psi a}^t) = v)$$

Also, for notational convenience, and without loss of generality, we express our prior beliefs in terms of a set of observations, Ω_0 , at $t = 0$, so that $P_0(\sum_a T_{\psi a}^0) \propto \prod_{(\omega, v) \in \Omega_0} P(\omega(T_{\psi a}^0) = v)$ (this also allows any priors that aren't at $t = 0$). The posterior can now be written as

$$P(T_{\psi a}^t | \Omega) \propto \begin{cases} \prod_{(\omega, v) \in \Omega} P(\omega(T_{\psi a}^t) = v) \prod_{t, \psi, a} P(\pi(\psi, \sum_b T_{\phi b}^t) = a)^{T_{\psi a}^t} & \text{if } T_{\psi a}^t \text{ is valid} \\ 0 & \text{otherwise} \end{cases} \quad (3)$$

where Ω now includes any priors.

In many practical applications, sampling from this posterior is difficult because it has zero probability for the vast majority of trajectories (i.e. most trajectories are invalid, contain an impossible action or are refuted by the observations). This is true of our predator/prey model, for example. Even though we can generate valid trajectories that fit the prior by simply performing a forward execution of the model from an initial state drawn from the prior, if the trajectory has no predator or prey at a time and place that we observed footprints, then the observations refute the trajectory and the probability falls to zero. It doesn't take many observations until the probability of randomly choosing a trajectory that fits the observations becomes very small indeed. So simple techniques such as rejection sampling are not practical.

In this paper we'll use the Metropolis-Hastings algorithm to generate samples. However, this method requires a proposal function which randomly generates the next sample given the current one. However, given the current sample, we still have the problem of finding a trajectory with non-zero probability. For example if we generate a new trajectory by perturbing one or more elements of the last sample at random (a common strategy with Metropolis-Hastings), it's very unlikely that we'll end up with a trajectory that is valid, contains only possible actions and satisfies the observations. So, the proposed next sample would almost certainly be rejected and we'd probably end up stuck on the first sample until we grew old.

2 Approximating the support of the posterior

To solve this problem we'd like to constrain our proposal function to the support of the posterior, $\text{supp}(P(T_{\psi a}^t | \Omega))$ (i.e. the set of trajectories that have non-zero probability).

If we let

$$\pi'(s, \xi, T_{\psi a}^t) = \pi(\xi, \sum_b T_{\phi b}^s)$$

Then from equation 3

$$\text{supp}(P(\cdot | \Omega)) = \bigcap_{(\omega, v) \in \Omega} \text{supp}(P(\omega(\cdot) = v)) \cap \bigcap_{t, \psi, a} \text{supp}(P(\pi'(t, \psi, \cdot) = a)) \cap \{T_{\psi a}^t \mid T_{\psi a}^t \text{ is valid}\} \quad (4)$$

i.e. in order for the posterior to be non-zero, all the observation likelihoods must be non-zero, the probability of each agent's action at every timestep must be non-zero and the trajectory must be valid.

The first two terms in equation 4 consists of the supports of computer programs whose inputs are ABM trajectories and whose outputs are given, i.e. the set of trajectories that, when passed to a computer program, would produce a given output.

Calculating the support of a computer program for a given output is, in full generality, NP-complete² but it is possible to use a technique known as *abstract interpretation* (Cousot & Cousot, 1977) to efficiently calculate a superset of the support. So, given a computer program ρ , we can calculate a set $\mathcal{P}(\rho, v)$ such that

$$\text{supp}(P(\rho(\cdot) = y)) \subset \mathcal{P}(\rho, v)$$

Tools to perform abstract interpretation already exist (e.g. PAGAI (Henry, Monniaux, & Moy, 2012)) and are used widely in applications such as the verification of safety critical systems (Blanchet et al., 2003) and in practice $\mathcal{P}(\rho, v)$ is often reasonably tight (i.e. most members of $\mathcal{P}(\rho, v)$ are in $\text{supp}(P(\rho(\cdot) = y))$). For our application we choose to express $\mathcal{P}(\rho, v)$ in terms of a set of linear inequalities on ρ 's inputs, this corresponds to the abstract domain of convex polyhedra (Cousot & Halbwachs, 1978) (Becchi & Zaffanella, 2018). Calls to the random number generator can be dealt with in the abstract domain by generating a new variable, r , that satisfies $0 \leq r < 1$ for each call to $\text{Random}()$.

If we replace the supports in equation 4 with their equivalent supersets we end up with a superset of the posterior

$$\text{supp}(P(\cdot|\Omega)) \subset \bigcap_{(\omega, v) \in \Omega} \mathcal{P}(\omega, v) \cap \bigcap_{t, \psi, a} \mathcal{P}(\pi'(t, \psi, \cdot), a) \cap \{T_{\psi a}^t \mid T_{\psi a}^t \text{ is valid}\} \quad (5)$$

The intersections are now intersections of sets of inequalities which are just the set of trajectories that satisfy all of the inequalities, i.e. the union of the inequalities.

Constraining our proposal function to members of the superset in equation 5 instead of the true support won't affect the stationary distribution of the Markov Chain. If the proposal function happens to return a trajectory that isn't in $\text{supp}(P(\rho(\cdot) = y))$ then it will just be rejected. This is fine as long as we generate acceptable proposals at a reasonable rate.

It remains for us to deal with the final term in equation 5, which requires the trajectories in the support to be valid. The first part of validity, the continuity constraints in 2, are not a problem because they are also linear constraints so can be treated along with the other constraints. The non-negative integer constraints of equation 1 are left as separate constraints.

Without loss of generality, we express inequalities as equalities by introducing slack variables so that

$$\sum_i c_i x_i \leq y$$

is equivalent to

$$\begin{aligned} \sum_i c_i x_i + \lambda &= y \\ \text{subject to } \lambda &\geq 0 \end{aligned}$$

All the constraints can now be gathered into a standard matrix form

$$\begin{aligned} AX &= B \\ \text{subject to } x_i &\geq 0 \\ x_i &\in \mathbb{Z} \end{aligned} \quad (6)$$

where the vector x contains the elements of the trajectory tensor $T_{\psi a}^t$ and the slack variables λ .

3 From convex polyhedron to Markov process

Given a set of inequalities that contain the support of the posterior, we next need to define a stochastic proposal function $f : \mathbb{S} \rightarrow \mathbb{S}$ on a set of Markov states, \mathbb{S} , where each state, σ , is associated with a trajectory $T(\sigma)$. The proposal function should have the following properties:

- There should be a set of transitions with non-zero probability between any two Markov states. This ensures proper mixing as time tends to infinity.

²Consider, for example, a program that accepts an assignment of variables to truth values, and returns true if that assignment satisfies a Boolean formula. Deciding whether the support of this program, given that it returns true, is empty or not is equivalent to solving the Boolean satisfiability problem, which is known to be NP-complete (Cook, 1971)

- The probability of being in a state associated with trajectory T should be our target posterior probability $P(T | \Omega)$.
- For any transition from state $s_a \rightarrow s_b$ the probability of transitioning from $s_b \rightarrow s_a$ should be non-zero. This allows us to attain detailed balance by using the Metropolis Hastings algorithm.
- The proposal function should be computationally efficient.

3.1 A proposal function for a Fermionic ABM

We define a Fermionic ABM to be one where no two agents can share the same state at the same time. i.e. in addition to the validity constraints of equations 1 and 2 the trajectory of a Fermionic ABM also satisfies the *Fermionic constraint*

$$\forall t, \psi : \sum_a T_{\psi a}^t \leq 1 \quad (7)$$

in which case we call it a *Fermionic trajectory* [this can be weakened to $T_{\psi a}^t \leq 1$].

Theorem 1. *All trajectories that satisfy equations 1 and 7 are at extreme points. Where an extreme point is one that can't be expressed as a linear combination of two other points that satisfy the constraints.*

Proof. Inequality 7 can only be satisfied by non-negative integers if, for any given t and ψ either all $T_{\psi a}^t$ are zero or one element of $T_{\psi a}^t$ is one and the rest zero. In the first case the solution is touching all $|a|$ of the $T_{\psi a}^t \geq 0$ constraints, in the second case the solution is touching $|a| - 1$ of the $T_{\psi a}^t \geq 0$ constraints and one $\sum_a T_{\psi a}^t \leq 1$ constraint to give a total of $|a|$.

[Conversely, we can show that all extreme points are integer points, so we have the polytope of integer solutions]

[In the case of the action Fermionic constraint, this can be seen immediately since equations 1 and the action Fermionic constraint describe a unit hypercube, so all integer solutions are on vertices.]

□

This seemingly abstract fact has a profound practical consequence for the efficiency of our proposal function. Given a system of m linear constraints on n variables expressed in the form $AX = B$, let a *pivot state* be a partition of the variables into m *basic variables* and $n - m$ *non-basic variables* so we can express the constraints in the form

$$A_b X_b + A_n X_n = B \quad (8)$$

where X_b and X_n are the basic and non-basic variables respectively, A_b is the matrix formed from the columns in A that correspond to the basic variables and A_n is the matrix formed from the columns of A that correspond to the non-basic variables. If we now constrain all non-basic variables to zero then we're left with $A_b X_b = B$ but since A_b is an $m \times m$ square matrix then $X_b = A_b^{-1} B$ gives a unique solution, as long as A_b is invertable. If A_b is invertable and X_b has no negative entries then we say that the pivot state is valid. It can be shown that all valid pivot states of equations 6 correspond to extreme points (Dantzig, Orden, Wolfe, et al., 1955).

This idea leads to a computationally very efficient way of perturbing trajectories. We begin by representing the constraints and the current trajectory in the form of equation 8. If we multiply any row of A_b by some constant, c , then the truth of equation 8 can be maintained by multiplying the same row of A_n and of B by c as well. Similarly if we add one row of A_b to another row, the equations can be maintained by performing the same row operation on A_n and B . If the current trajectory is valid, then A_b must be invertable, so it is possible, by performing a sequence of row additions and multiplications, to transform A_b into a form such that each row and column is zero on all elements except one, which is 1. We'll call this the standard form. Once in this form, the solution X_b can be read off directly from B if we assume $X_n = 0$.

Given a valid solution in standard form, a perturbed valid solution can be generated by choosing a column, C , of A_n , with at least one positive-valued row, and choosing from among those rows the one, C_i , that minimises B_i/C_i . We then make the variable corresponding to column C a basic variable and make the basic variable corresponding to row i of A_b non-basic by swapping column C with the column on row i of A_b that has value 1. A_b is now not in standard form as column C will, in general, have more than one non-zero value so we perform row additions and multiplications to return it to standard form. This perturbation is known as a *pivot operation* and is the same method as employed in the Simplex algorithm (Dantzig et al., 1955) (Vanderbei, 2020). This is exactly what we need for our proposal function since it efficiently generates a new valid trajectory from a current one.

Although pivoting ensures solutions are positive, there may exist extreme points that don't fall on the grid of integer solutions and so do not satisfy the integer constraint. However, if we posit the existence of a *null trajectory* and associate all fractional solutions with the null trajectory, then simply ignore null samples, then we end up with samples from the posterior, as required (i.e. the Markov chain can pass through fractional states, but no sample is generated when doing so). However, we don't want to spend too much computational effort moving between fractional states, so we encourage

the Markov process away from these fractional states by multiplying their probability by a “fractional penalty” e^{-kn} where n is the number of non-integer values in the solution and k is a constant parameter.

The use of pivoting is complicated somewhat when the solutions are degenerate. A degenerate solution is one that can be generated by more than one pivot state. This can occur when a pivot state has basic variables with value zero (in which case we say the pivot state is degenerate and call the zero variables *degenerate variables* of the pivot state). Pivoting on a degenerate variable does not change the solution since during a pivot the outgoing basic variable becomes constrained to zero; but if the basic-variable’s value was zero to start with, constraining it to zero has no effect on the solution (for a thorough theoretical development of degeneracy see (Zörnig, 1993)). This is problematic because we need to assign a probability to each Markov state but under degeneracy a pivot state can’t be given the probability of its trajectory because this would result in degenerate trajectories being “counted” multiple times, once for each degenerate pivot state of the trajectory. To maintain the correct distribution, the probabilities of all the degenerate pivot states of a given trajectory should sum to the probability of that trajectory.

To solve this problem in a computationally efficient way we introduce the concept of an *ordered pivot state* which consists of a partition of the variables into basic and non-basic variables as before, but in addition there is now an ordering of the basic variables. Since we represent A_b as a matrix in our standard form representation, it is easy to represent this ordering in the ordering of the rows of A_b . As before, the validity of equation 8 can be maintained under permutations of the ordering of the rows of A_b by performing the same row permutations on A_n and B . We now associate a Markov state with each ordered pivot state, assigning it a probability according to algorithm 3:

Algorithm 3 Algorithm to calculate probability of an ordered pivot state

```

function PROBABILITY( $A_b, A_n, B$ ) ▷ Standard representation of an ordered pivot state
   $P_T \leftarrow$  the probability of the trajectory of the current solution
   $P_d \leftarrow 1.0$ 
   $R \leftarrow$  indices of  $B$  that are zero, in descending order
   $\sigma \leftarrow$  length of  $R$ 
   $C \leftarrow \emptyset$ 
  for  $i = 0 \dots \sigma - 1$  do
     $C \leftarrow C \cup$  the set of non-zero columns in row  $R[i]$  of  $A_n$ 
     $P_d \leftarrow P_d / (|C| + i + 1)$ 
  end for
  return  $\frac{P_d P_T (|B| - \sigma)!}{|B|!}$ 
end function

```

The proposal function for ordered pivot states is shown in algorithm ??

Algorithm 4 Proposal function for ordered pivot states

```

function PROPOSAL( $A_b, A_n, B$ ) ▷ Standard representation of the current ordered pivot state
   $A'_b, A'_n, B' \leftarrow A_b, A_n, B$ 
   $\alpha, \beta \leftarrow$  constant parameters
   $p_0 \leftarrow$  the set of degenerate pivots of  $(A_b, A_n, B)$  (i.e. the ones that do not change the solution)
   $p_1 \leftarrow$  the set of non-degenerate pivots of  $(A_b, A_n, B)$  (i.e. the ones that change the solution)
   $P_{f/b} \leftarrow 1$  ▷ the ratio of probabilities of making this transition forwards/backwards
  if BERNOULLI( $\alpha$ ) then
     $r \leftarrow$  choose member of  $p_1$  with uniform probability
    perform pivot  $r$  on  $(A'_b, A'_n, B')$ 
     $p'_0 \leftarrow$  number of non-degenerate pivots of  $(A'_b, A'_n, B')$ 
     $P_{f/b} \leftarrow \frac{|p'_1|}{|p_1|}$ 
  else if BERNOULLI( $\beta$ ) then
     $r \leftarrow$  choose member of  $p_0$  with uniform probability
    perform pivot  $r$  on  $(A'_b, A'_n, B')$ 
     $p'_0 \leftarrow$  number of degenerate pivots of  $(A'_b, A'_n, B')$ 
     $P_{f/b} \leftarrow \frac{|p'_0|}{|p_0|}$ 
  else
    choose two rows  $a$  and  $b$  with uniform probability
    swap rows  $a$  and  $b$  of  $(A'_b, A'_n, B')$ 
  end if
  return  $(A'_b, A'_n, B', P_{f/b})$ 
end function

```

We now show that this assignment of probability to ordered pivot states has the property that the sum of the probabilities associated with a trajecotry is the probability of that trajectory.

First some notation: let $\mathbb{E}(T)$ be the set of all ordered pivot states with solution T , let $T(S)$ be the trajectory associated with pivot state S , let $D(S)$ be the tuple of degenerate variables in ordered pivot state S , let $D_{\leq n}(S)$ be the n -tuple consisting of the first n elements of $D(S)$ and let

$$C(\mathbb{S} | V) = \left\| \left\{ V^+ : S \in \mathbb{S}, V^+ = D_{\leq |V|+1}(S), V_{\leq |V|}^+ = V \right\} \right\| \quad (9)$$

i.e. among the members of \mathbb{S} that have the degenerate variable prefix V , $C(\mathbb{S}|V)$ counts how many distinct degenerate variable prefixes of length $|V| + 1$ there are.

Theorem 2. *If we assign the probability*

$$P(S) = \frac{(m - \sigma_{T(S)})!}{m!} \frac{P_{T(S)}}{\prod_{q=1}^{\sigma_{T(S)}} C(\mathbb{E}(T(S)) | D_{\leq (q-1)}(S))}$$

to pivot state S , where m is the number of constraints and σ_T is the degeneracy of trajectory T (i.e. m minus the number of non-zero variables in T) then

$$\sum_{S \in \mathbb{E}(T)} P(S) = P_T$$

Proof. Expanding the sum over degenerate states

$$\sum_{S \in \mathbb{E}(T)} P(S) = P_T \frac{(m - \sigma_T)!}{m!} \sum_{S \in \mathbb{E}(T)} \frac{1}{\prod_{q=1}^{\sigma_T} C(\mathbb{E}(T) | D_{\leq (q-1)}(S))}$$

Notice first that all pivot states with the same σ_T -tuple of degenerate variables, $D(S)$, have the same probability. There are exactly $\frac{m!}{(m - \sigma_T)!}$ ordered pivot states with a given $D(S)$, corresponding to the different placings of the $m - \sigma_T$ non-degenerate basic variables among the m possible positions. So, if we let

$$\mathbb{D}(T) = \{D(S) : S \in \mathbb{E}(T)\}$$

then

$$\sum_{S \in \mathbb{E}(T)} P(S) = P_T \sum_{D \in \mathbb{D}(T)} \frac{1}{\prod_{q=0}^{\sigma_T-1} C(\mathbb{E}(T) | D_{\leq q})}$$

Since $C(\mathbb{S} | V)$ depends only on the order of the degenerate variables in the members of \mathbb{S}

$$C(\mathbb{E}(T) | V) = C(\mathbb{D}(T) | V)$$

so

$$\sum_{S \in \mathbb{E}(T)} P(S) = P_T \sum_{D \in \mathbb{D}(T)} \frac{1}{\prod_{q=0}^{\sigma_T-1} C(\mathbb{D}(T) | D_{\leq q})} \quad (10)$$

if we separate the terms in the sum into groups that share the same prefix apart from the last variable then we get

$$\sum_{S \in \mathbb{E}(T)} P(S) = P_T \sum_{V : D \in \mathbb{D}(T), V = D_{\leq (\sigma_T-1)}} \left(\frac{1}{\prod_{q=0}^{\sigma_T-2} C(\mathbb{D}(T|V_{\leq q}))} \sum_{V^+ : V^+ \in \mathbb{D}(T), V = V_{\leq (\sigma_T-1)}^+} \frac{1}{C(\mathbb{D}(T|V^+))} \right)$$

So

$$\sum_{S \in \mathbb{E}(T)} P(S) = P_T \sum_{V : D \in \mathbb{D}(T), V = D_{\leq (\sigma_T-1)}} \frac{1}{\prod_{q=0}^{\sigma_T-2} C(\mathbb{D}(T|V_{\leq q}))} \frac{\left\| \left\{ V^+ : V^+ \in \mathbb{D}(T), V = V_{\leq (\sigma_T-1)}^+ \right\} \right\|}{C(\mathbb{D}(T|V_{\leq (\sigma_T-1)}))}$$

but the final fraction here is 1 by equation 9, so

$$\sum_{S \in \mathbb{E}(T)} P(S) = P_T \sum_{V : D \in \mathbb{D}(T), V = D_{\leq (\sigma_T-1)}} \frac{1}{\prod_{q=0}^{\sigma_T-2} C(\mathbb{D}(T|V_{\leq q}))}$$

we can now use the same argument again to reduce the upper bound of q iteratively until we show that the sum on the right hand equals one, so

$$\sum_{S \in \mathbb{E}(T)} P(S) = P_T$$

□

Theorem 3. *Algorithm 3 assigns a probability*

$$P(S) = \frac{(m - \sigma_{T(S)})!}{m!} \frac{P_{T(S)}}{\prod_{q=1}^{\sigma_{T(S)}} C(\mathbb{E}(T(S)) \mid D_{\leq(q-1)}(S))}$$

to pivot state S

Proof.

□

Given this, it's clear that all the necessary properties of the proposal function are fulfilled.

References

- Becchi, A., & Zaffanella, E. (2018). An efficient abstract domain for not necessarily closed polyhedra. In *International static analysis symposium* (pp. 146–165).
- Blanchet, B., Cousot, P., Cousot, R., Feret, J., Mauborgne, L., Miné, A., . . . Rival, X. (2003). A static analyzer for large safety-critical software. In *Proceedings of the acm sigplan 2003 conference on programming language design and implementation* (pp. 196–207).
- Cook, S. A. (1971). The complexity of theorem-proving procedures. In *Proceedings of the third annual acm symposium on theory of computing* (pp. 151–158).
- Cousot, P., & Cousot, R. (1977). Abstract interpretation: a unified lattice model for static analysis of programs by construction or approximation of fixpoints. In *Proceedings of the 4th acm sigact-sigplan symposium on principles of programming languages* (pp. 238–252).
- Cousot, P., & Halbwachs, N. (1978). Automatic discovery of linear restraints among variables of a program. In *Proceedings of the 5th acm sigact-sigplan symposium on principles of programming languages* (pp. 84–96).
- Dantzig, G. B., Orden, A., Wolfe, P., et al. (1955). The generalized simplex method for minimizing a linear form under linear inequality restraints. *Pacific Journal of Mathematics*, 5(2), 183–195.
- Gal, T., & Geue, F. (1992). A new pivoting rule for solving various degeneracy problems. *Operations Research Letters*, 11(1), 23–32.
- Henry, J., Monniaux, D., & Moy, M. (2012). Pagai: A path sensitive static analyser. *Electronic Notes in Theoretical Computer Science*, 289, 15–25.
- Vanderbei, R. J. (2020). *Linear programming* (Vol. 5). Springer.
- Yamada, T., Yoruzuya, J., & Kataoka, S. (1994). Enumerating extreme points of a highly degenerate polytope. *Computers & operations research*, 21(4), 397–410.
- Zörnig, P. (1993). A theory of degeneracy graphs. *Ann Oper Res*, 46, 541–556. doi: 10.1007/BF02023113